


Introduction to Amazon EKS Auto Mode

Mohan Ramadoss



kubernetes




Amazon EKS

What Is Amazon EKS Auto Mode?

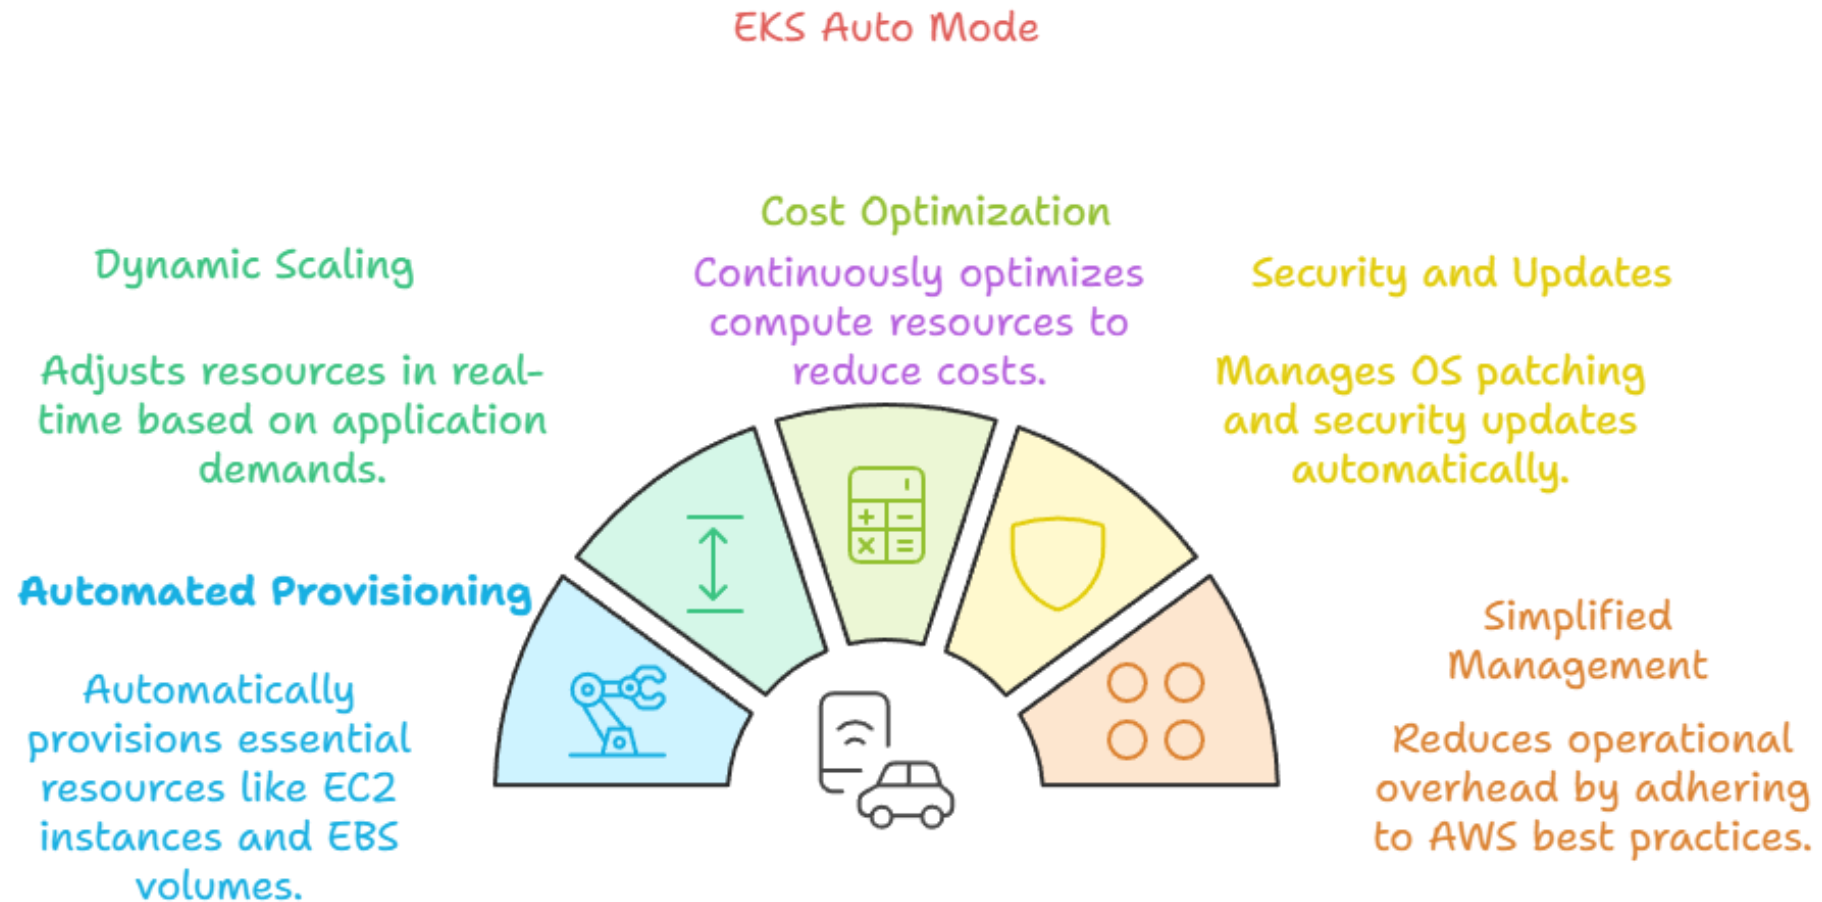
Automated Management in EKS Auto Mode

Kubernetes Cluster Management



Optimized Kubernetes Environment

Key Features of EKS Auto Mode






Amazon EKS cluster



EKS account

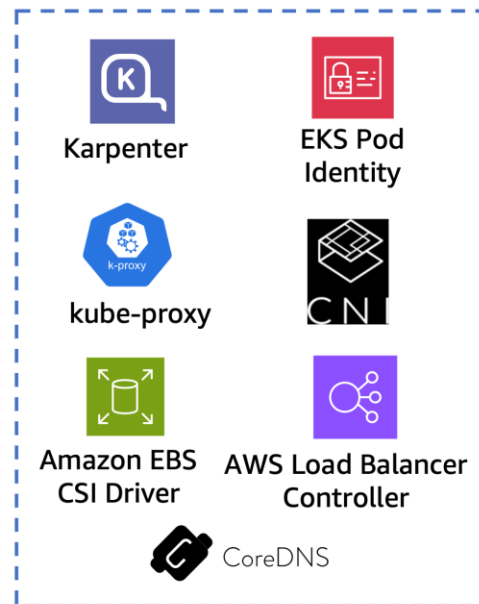


API Server
Instances




etcd
Instances

Managed Control Plane



Amazon EKS Add-Ons



EC2 Instances

Customer account



Application Load
Balancer



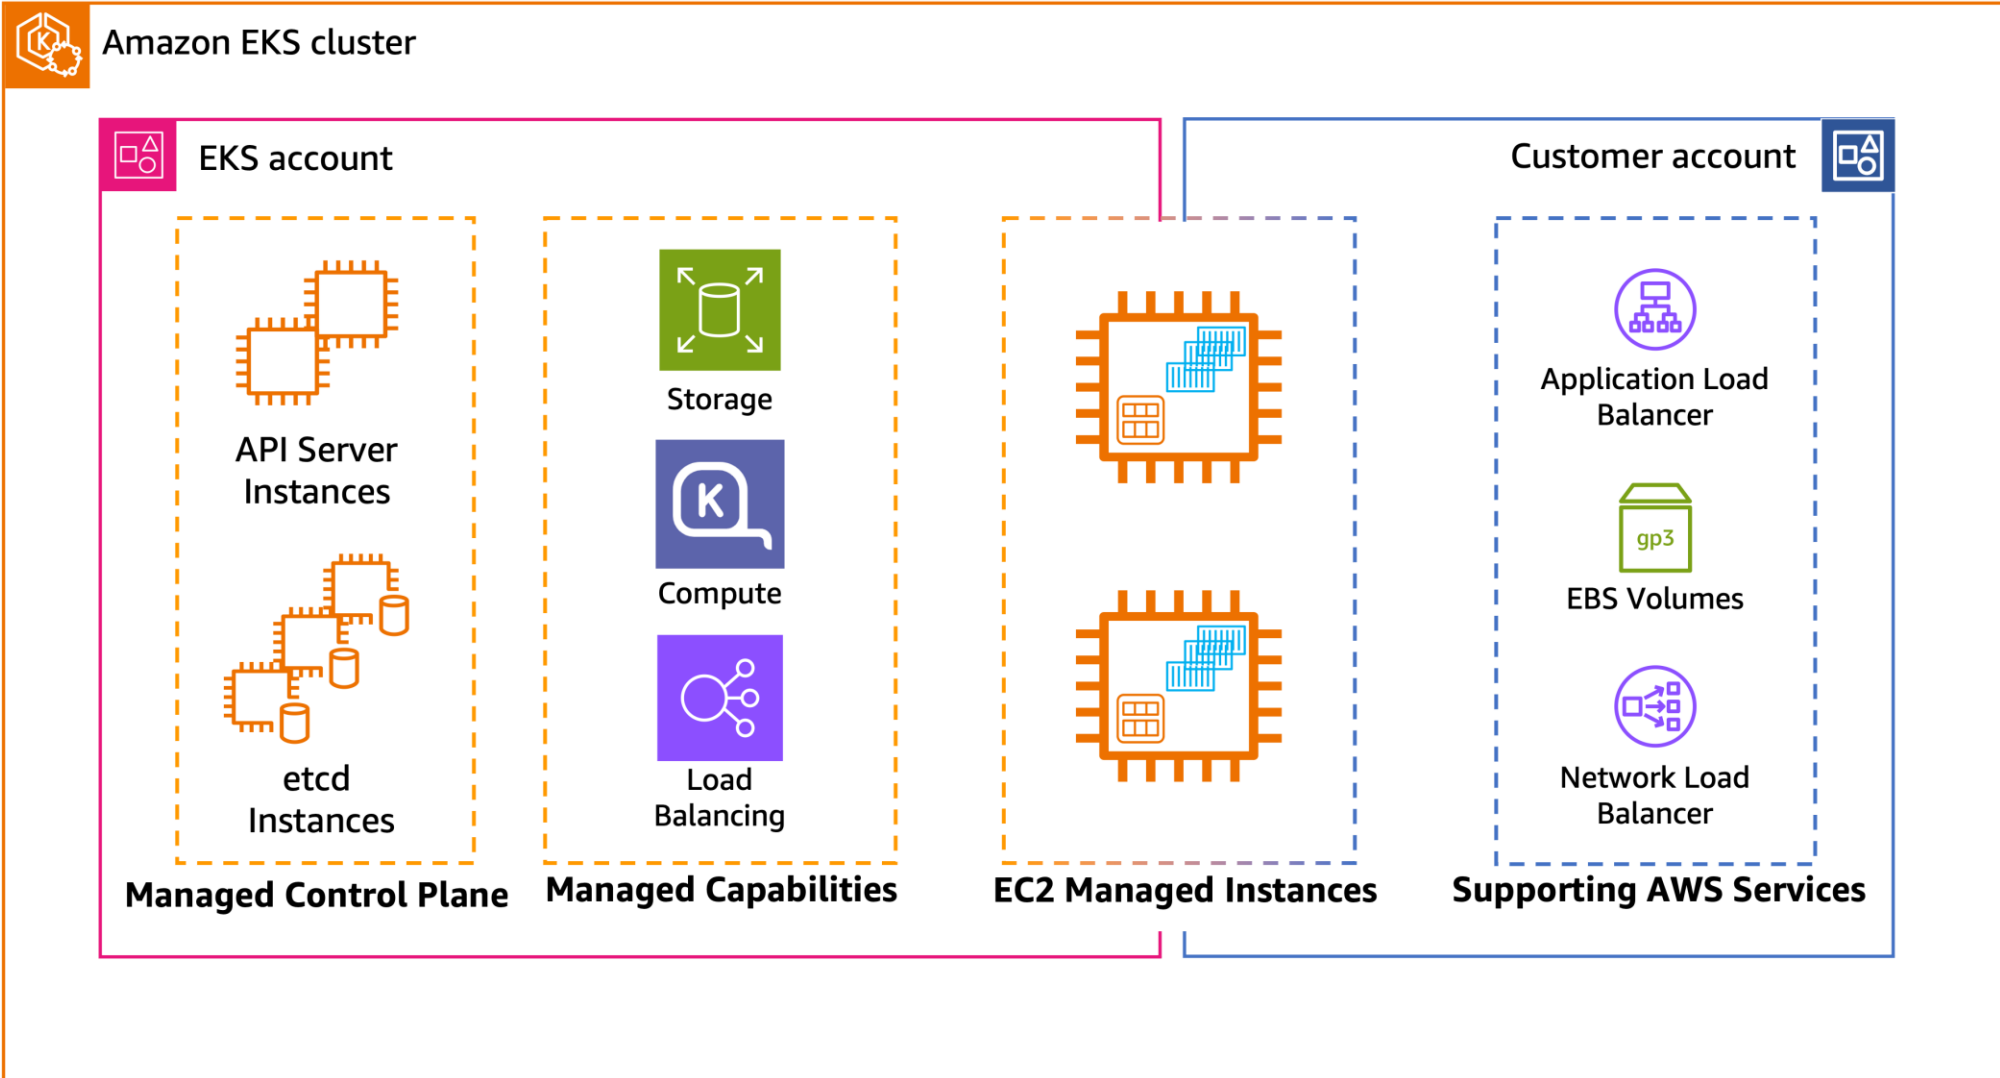
EBS Volumes



Network Load
Balancer

Supporting AWS Services

Before Auto Mode



After Auto Mode

Shared Responsibility Model with EKS previously

System Agents,
Kubelet,
Container Runtime

Application containers: Availability, security, monitoring

VPC infrastructure, cluster configuration, add-ons

Cluster EC2
Instances

Lifecycle

Operating
System

Patching

Monitoring

Health
and Repair

Cluster
Capabilities

Compute
Autoscaling

Pod Networking
/ Network Policy

Elastic Load
Balancing

Storage Drivers

EKS Cluster
Control Plane

KBs API Server

X-ENIs

KBs Etcd Database

Foundation
Services

Compute

Storage

Networking

Monitoring

AWS Global
Infrastructure

Regions

Local Zones

Edge Locations



Managed by
Customers



Managed by
Amazon
Web Services

Shared Responsibility Model with EKS Auto Mode

System Agents,
Kubelet,
Container Runtime

Application containers: Availability, security, monitoring

VPC infrastructure, cluster configuration, add-ons

Customer
IAM

Managed by
Customers

Cluster EC2
Instances

Lifecycle

Operating
System

Patching

Monitoring

Health
and Repair

Cluster
Capabilities

Compute
Autoscaling

Pod Networking
/ Network Policy

Elastic Load
Balancing

Storage Drivers

EKS Cluster
Control Plane

K8s API Server

X-ENIs

K8s Etcd Database


Managed by
Amazon
Web Services

Foundation
Services

Compute

Storage

Networking

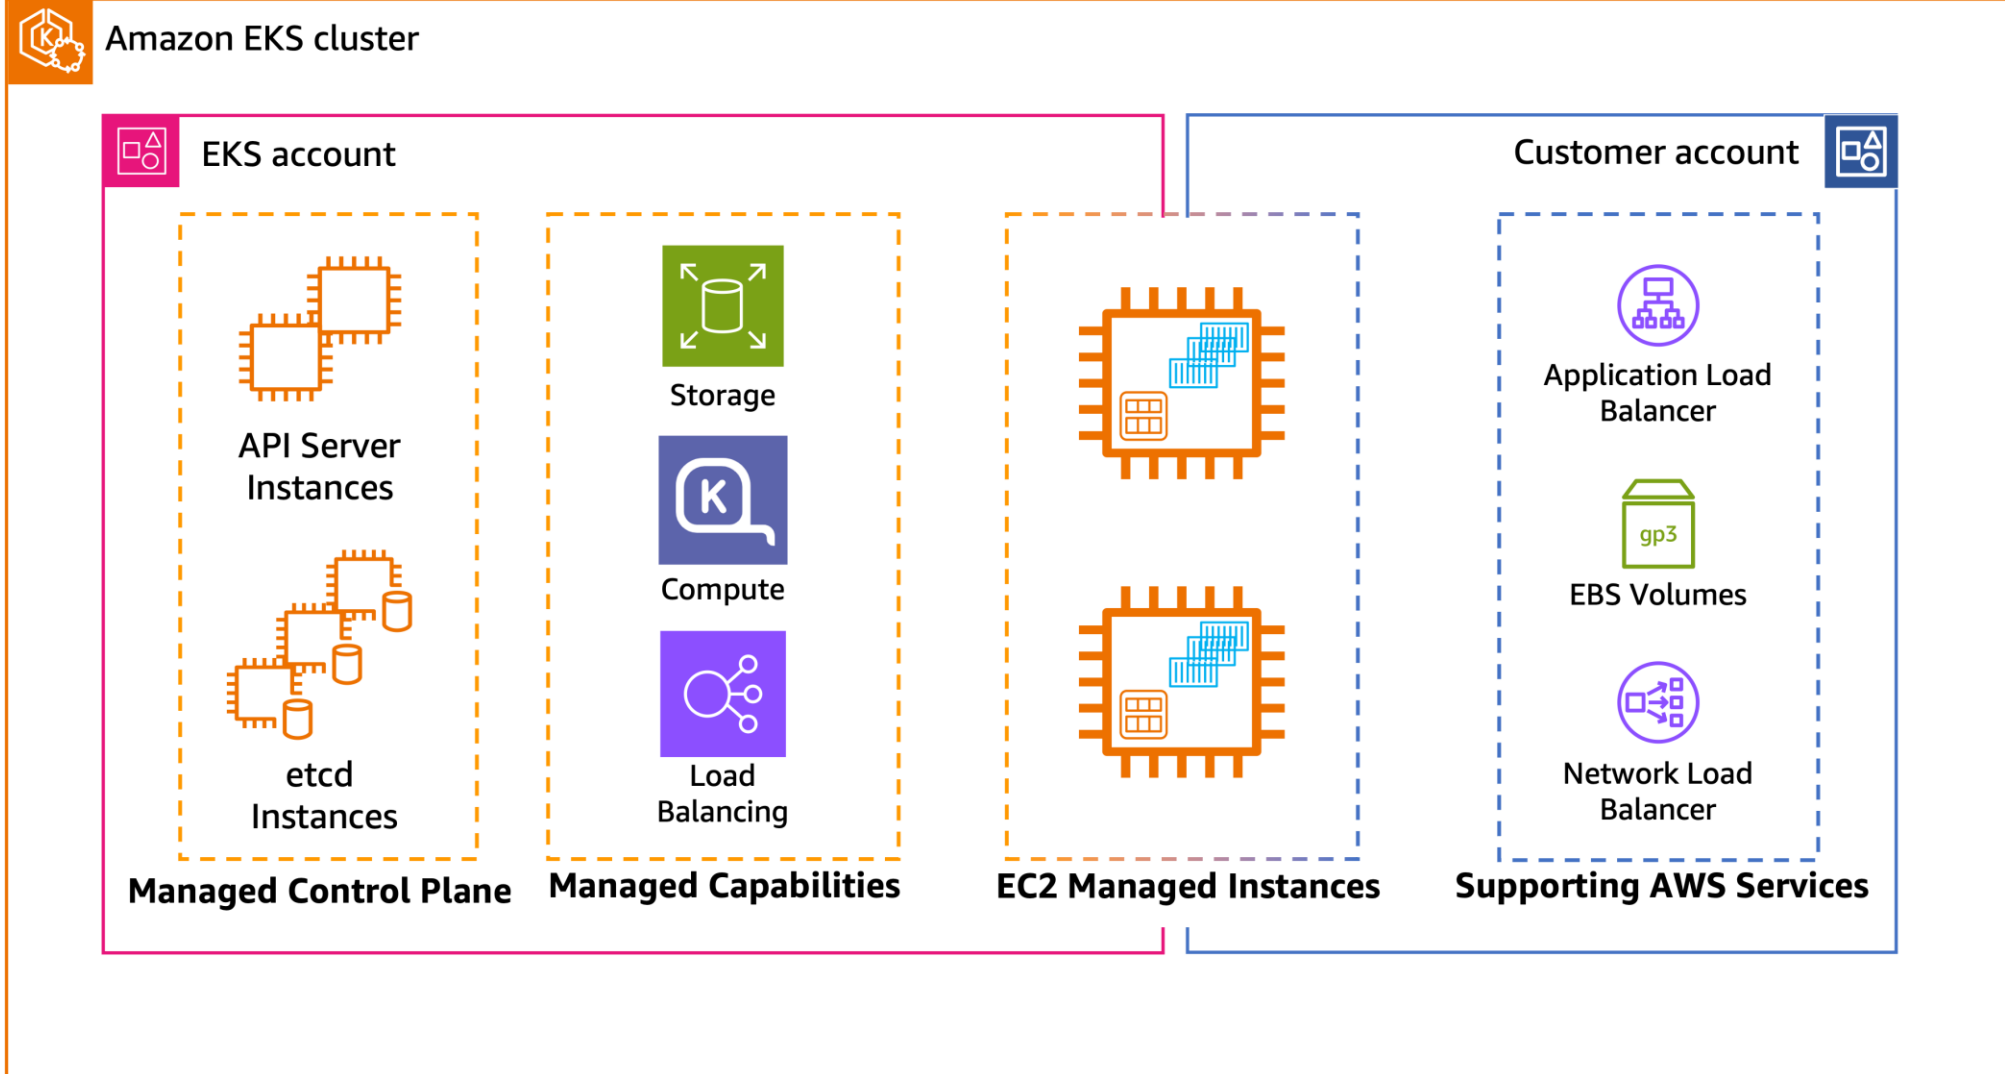
Monitoring

AWS Global
Infrastructure

Regions

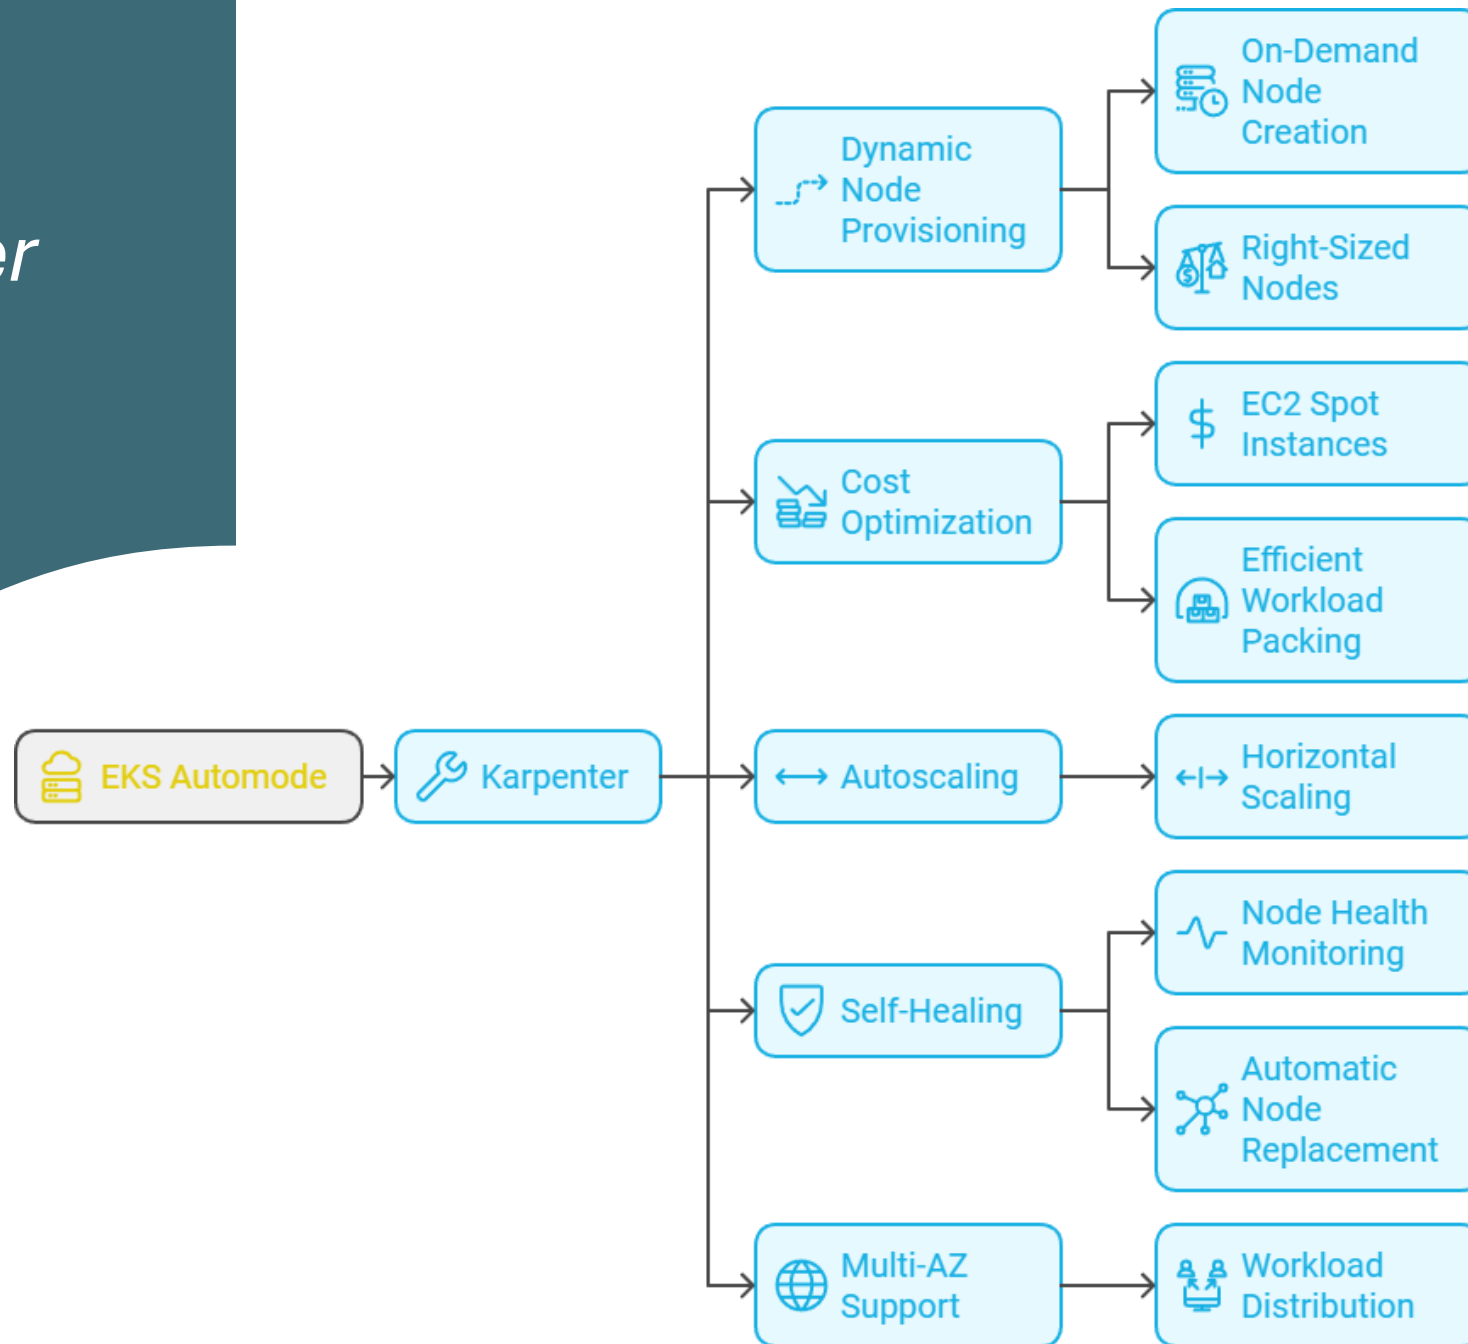
Local Zones

Edge Locations



Architecture overview

Karpenter




Workflow of EKS Auto Mode with Karpenter



Key Functionalities of EKS Auto Mode

- **Node Lifecycle Management:** Manages the lifecycle of EC2 instances, including scaling and replacement of unhealthy nodes.
- **System Processes:** Runs node capabilities as system processes managed by AWS, not as Kubernetes DaemonSets.
- **Automatic Updates:** Handles cluster upgrades and OS updates automatically.


Understanding EKS Auto Mode Functionalities



Advanced Use Cases and Configurations

- **High Availability:** Configure pod topology spread constraints for high availability across nodes and zones.
- **Custom Compute Selection:** Use Kubernetes labels to specify compute requirements for workloads.
- **Application Best Practices:** Implement pod disruption budgets, resource requests/limits, and graceful shutdowns.

Enhancing EKS Auto Mode with Advanced Strategies



Method	Description	Use Case
AWS Management Console	Simple graphical interface to configure and create the cluster with Auto Mode in a few clicks.	Quick setups and minimal technical expertise.
AWS CLI	Command-line tool to script and automate the creation of Auto Mode clusters.	Scripted workflows and automation.
Terraform	Infrastructure as Code (IaC) tool for declarative provisioning of EKS clusters.	Scalable and repeatable setups.
AWS CloudFormation	Service to define and deploy infrastructure templates, including EKS Auto Mode clusters.	Integrations with broader AWS setups.
CDK (AWS Cloud Development Kit)	Use programming languages like Python, JavaScript, or TypeScript to define cluster infrastructure.	Developers preferring code-driven approaches.

DEMO

- **Deploy a sample inflate workload to an Amazon EKS Auto Mode cluster**

Shows how to deploy a sample workload to an EKS Auto Mode cluster using kubectl commands.

- **Deploy a sample load balancer workload to EKS Auto Mode**

Shows how to deploy a containerized version of the 2048 game on Amazon EKS.

- **Deploy a sample stateful workload to EKS Auto Mode**

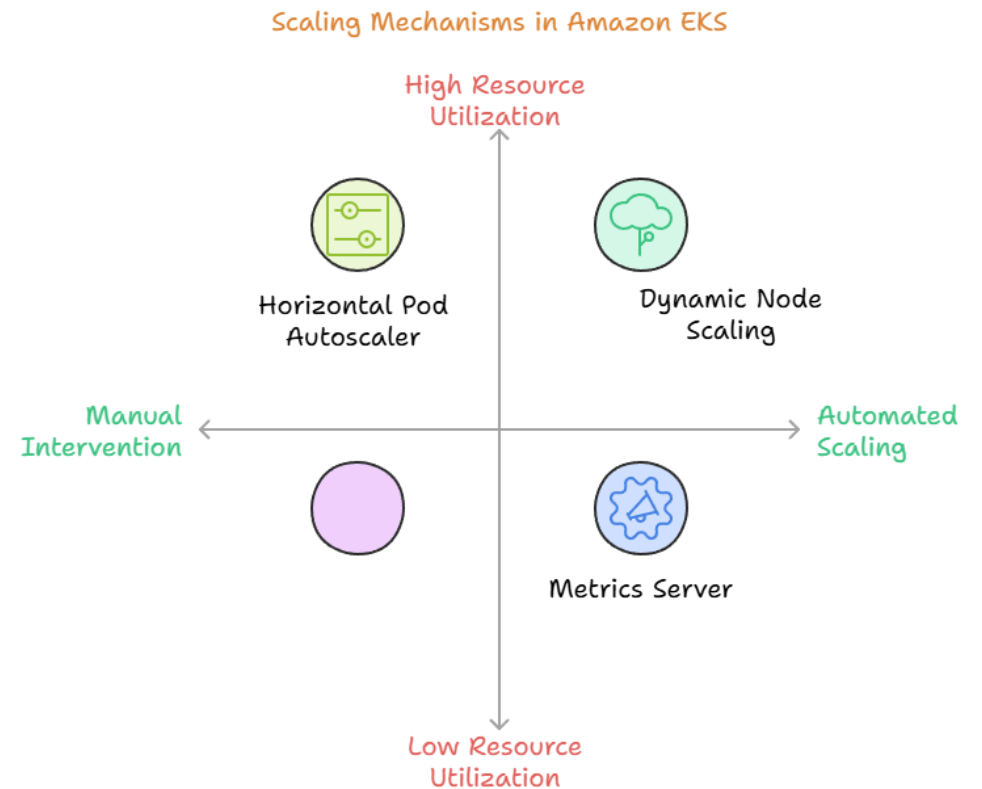
Shows how to deploy a sample stateful application to an EKS Auto Mode cluster.

Feature Comparison

Feature	EKS Auto Mode	EKS Standard Mode
Node Management	Fully managed - AWS manages the lifecycle of EC2 instances, including scaling, replacement of unhealthy nodes, and automatic updates.	User-managed - Users are responsible for managing EC2 instances, including scaling, updates, and node health.
Scaling	Automatic, optimized - Scales resources dynamically based on workload demands, ensuring efficient resource utilization.	Customizable - Users can configure scaling policies but must manage scaling themselves.
Target Audience	Beginners, teams wanting simplicity - Ideal for those who want to focus on application development rather than infrastructure management.	Advanced Kubernetes users - Suited for users who need fine-grained control over their Kubernetes environment.
Use Cases	Quick setups, auto-scaling - Perfect for development environments, batch processing, and containerized applications with fluctuating resource needs.	Fine-tuned, custom setups - Suitable for complex workloads requiring custom networking, advanced configurations, or specific compute requirements.
Operational Complexity	Reduced - Automates Kubernetes setup and infrastructure scaling, making it beginner-friendly.	Higher - Requires in-depth knowledge of Kubernetes and AWS services for optimal management.
Cost Efficiency	Optimized - Eliminates over-provisioning by scaling compute resources dynamically based on actual workload needs.	Manual optimization - Users must manually manage and optimize resource usage to control costs.
Security and Compliance	Improved - Managed updates, patching, and integration with AWS security services enhance cluster security.	User-managed - Users must ensure security updates and compliance themselves.
Flexibility	Limited - Provides a streamlined experience with less flexibility for custom configurations.	High - Offers full control over cluster configurations, networking, and compute selection.


Scaling Beyond Default Limits

- **Horizontal Pod Autoscaler (HPA):** Automatically scales the number of pods based on CPU usage or other metrics.
- **Metrics Server:** Collects resource usage data for HPA to make scaling decisions.
- **Dynamic Node Scaling:** EKS Auto Mode scales EC2 instances based on workload demands.



Benefits for Kubernetes Users

Unveiling the Benefits of EKS Auto Mode



Conclusion

